

WHERE IS HISTORY BEING WRITTEN?
GEOREFERENCING CONTRIBUTIONS
TO WIKIPEDIA

DAVID KALTSCHMIDT

Diplomarbeit

Dr. Claudia Müller-Birn
Prof. Dr. Robert Tolksdorf
Institut für Informatik
Freie Universität Berlin

David Kaltschmidt: *Where is history being written? Georeferencing contributions to Wikipedia*

Diplomarbeit, © August 2011 – February 2012

SUPERVISORS:

Dr. Claudia Müller-Birn

Prof. Dr. Robert Tolksdorf

LOCATION:

Berlin, Germany

YEAR:

August 2011 – February 2012

PRODUCED WITH:

L^AT_EX using the ClassicThesis package.

ABSTRACT

rewrite once all is done

Wikipedia is more than an online encyclopedia. It is also a news channel as well as a self-updating history book. A global readership can follow political events as they unfold, written about by local people and later edited by other volunteers. This thesis describes a method to answer the question to what extent local volunteers write about events in their own country. First, the geographic origin of each individual article contribution is determined. In a second step, a given article is annotated with georeferences on a word level. The properties of these annotations then allow for a statistical geographic analysis of a single article or a category of articles.

ZUSAMMENFASSUNG

rewrite once all is done

Als Online-Enzyklopädie ist Wikipedia nicht nur Nachschlagewerk sondern auch ein sich stetig wandelndes Geschichtsbuch. Eine global verteilte Nutzerschaft liest und schreibt über lokale Ereignisse noch während sie passieren. Diese Arbeit beschreibt eine Methode zur Bestimmung des Anteils an Beiträgen, die vom betreffenden Land ausgehen. In einem ersten Schritt werden die geographischen Ursprünge aller Beiträge eines Artikels ermittelt. Mit den daraus erhaltenen Georeferenzen wird der Artikel Wort für Wort annotiert. Basierend auf diesen Annotationen kann dann der lokale Autoren-Anteil bestimmt werden.

CONTENTS

I	THOUGHTS	1
1	INTRODUCTION	2
1.1	Structure	3
2	FOUNDATION	5
2.1	Wikipedia	5
2.1.1	History	6
2.1.2	Wikimedia Foundation	6
2.1.3	Anatomy of an article	7
2.1.4	Categories	8
2.2	MediaWiki and editing	8
2.2.1	Templates	8
2.2.2	Revision history	9
2.2.3	Authors	10
2.2.4	User pages	10
2.3	Contributions	11
2.4	Georeferences	11
3	HYPOTHESES	13
3.1	Article creation	13
3.2	Participation	15
3.3	Text survival	16
II	METHODS	17
4	APPARATUS	18
4.1	Data sources	18
4.1.1	Wikipedia website	18
4.1.2	Database dumps	20
4.1.3	MediaWiki API	20
4.1.4	Toolserver	22
4.1.5	Third-party sources/Web services	23
4.2	Available Tools	25
4.2.1	Toolkits	25
4.2.2	Analysis projects	26
4.3	Application design	26
4.3.1	Technologies	26
4.3.2	Models	27
4.3.3	Views	28
4.3.4	Article analyzer	28
4.3.5	Group analyzer	29
4.4	Algorithms	29
4.4.1	Article requirements	30
4.4.2	Date parsing	30
4.4.3	Location parsing	32

4.4.4	Collective authorship	33
4.4.5	Locating users	33
4.4.6	Signature distance	35
4.4.7	Localness	38
4.5	Visualization	38
4.5.1	Maps	39
4.5.2	Line charts	40
4.5.3	Motion chart	41
4.6	Hypothesis analysis	41
4.7	Possible enhancements	45
4.7.1	Edit relevance	45
4.7.2	User page parsing	45
4.7.3	Geographic profiling	45
5	EXPERIMENTS	46
5.1	Data sets	46
5.1.1	By category	46
5.1.2	By template	46
5.1.3	Political article vs. place article	46
5.2	Application run	46
5.2.1	Distribution	46
5.2.2	Text survival	46
5.2.3	Localness	46
5.2.4	Motion chart	46
III	RESULTS	47
6	RESULTS	48
7	CONCLUSION	49
7.1	Limitations	49
7.1.1	Political events	49
7.1.2	Article location	50
7.1.3	Cross-language article growth	50
7.2	Further Research	50
IV	APPENDIX	51
	Bibliography	52

LIST OF FIGURES

Figure 1	The article <i>2011–2012 Bahraini uprising</i> viewed in a web browser on 00/00/0. 7
Figure 2	Revision history of the article <i>2011–2012 Bahraini uprising</i> on 00/00/0. 9
Figure 3	Article location for <i>2011–2012 Bahraini uprising</i> 39
Figure 4	Geographic origins by country for located authors of <i>2011–2012 Bahraini uprising</i> 39
Figure 5	Text survival in revision 471577075 grouped by country for located text of <i>2011–2012 Bahraini uprising</i> 40
Figure 6	Activity chart for <i>2011–2012 Bahraini uprising</i> 40
Figure 7	Temporal development of edit counts by country for <i>2011–2012 Bahraini uprising</i> 41
Figure 8	Temporal development of text proportion by country for <i>2011–2012 Bahraini uprising</i> 42

LIST OF TABLES

LISTINGS

1	Example JSON response to a query to list all bots that edited the article <i>2011–2012 Bahraini uprising</i> . . . 20
2	SoNet API response to a query for the article <i>2011–2012 Bahraini uprising</i> 22
3	Excerpt of the annotated markup for the revision 473029564 of the article <i>2011–2012 Bahraini uprising</i> 24
4	Date candidates algorithm 31
5	Date tokens 31
6	Article’s locate algorithm 32
7	User page location algorithm 34

8	Signature distance algorithm	37
9	Signature distance algorithm for all revisions	37
10	Localness of an author	38
11	Edit weight for map	40

ABBREVIATIONS

API	application programming interface
CSS	Cascading Style Sheets
CSV	comma-separated values
HTML	HyperText Markup Language
IP	Internet Protocol
ISO	International Organization for Standardization
JS	JavaScript
JSON	JavaScript Object Notation
NPOV	neutral point of view
PHP	PHP: Hypertext Preprocessor
XML	Extensible Markup Language

Part I

THOUGHTS

INTRODUCTION

If you are open to contributions from others, you generally end up with richer, better, more diverse and expert content than if you try to do it alone.¹

— Alan Rusbridger, editor of THE GUARDIAN

At the end of January 2011, when a wave of public protest spilled from Tunisia into Egypt, a small group of opposition parties and political activists called for a “Day of Rage” via Facebook, a social networking website. By January 25th their Facebook group had more than 80,000 supporters who drew attention to and helped organize the country-wide protests that followed. As people rallied the streets day after day, the Egyptian government first limited access to Twitter, a micro-blogging service, before cutting Egypt off the internet completely on January 28th.[2, 3]

In what came to be known as the Arab Spring, the use of online networks directly influenced the developing political situation. While Facebook played a part in organizing the protests, Twitter acted as an information channel during the demonstrations. As the events unravelled, they were reflected by articles created on Wikipedia, an online encyclopedia. Updated by the minute, the articles covering the protests formed a well of news reports.[4]

Wikipedia’s free access, open editing policy, and high quality level—putting it “head to head”[5] with Encyclopedia Britannica—turned it into a hugely popular website[6]. The server software used for the website, MediaWiki², ensures that the effort to change an article is minimal. Given an Internet connection and a web browser, anyone can add or edit an account of current events in a related article and publish it in a matter of seconds.

This form of news production turns the encyclopedia into a news channel that is constantly updated and corrected by an army of volunteers. The result is a self-governed news source that lends itself the aura of authority and credibility of a knowledge reference. At the same time, a technophile public that uses the Internet as an efficient means of news acquisition can check facts on Wikipedia, and act upon the consumed information.[7, p. 424–427] Therefore the collective authorship of such a news

¹ The Economist. *The people formerly known as the audience*. URL: <http://www.economist.com/node/18904124> (visited on 08/10/2011)

² <http://www.mediawiki.org> (visited on 10/31/2011)

medium could have a direct influence on the political decision-making process. As ordinary people become producers of journalism, the need arises to analyze these contributions. This thesis investigates how Wikipedia is being used during and after political events with a special focus on the geographic origins of contributions to Wikipedia articles treating those events.

Political events are often limited to a country or region. This is reflected by the Wikipedia articles covering the Arab Spring that will serve as examples throughout this thesis. There is an overarching parent article *Arab Spring*³ as well as single articles covering the revolution in each of the affected countries, e.g. Egypt⁴, Libya⁵, and Bahrain⁶. The events in Libya also exemplify how divided the political actors can be. While nearly all revolutionaries welcomed the airstrikes, one faction was concerned about foreign meddling and another one just opposed the deployment of ground troops.[8]

The collective authorship could be equally divided. Despite Wikipedia's core policy to oblige everyone to write from a neutral point of view (NPOV)⁷, people regularly express opinions. The collision of opinions in a collectively written article can result in a prolonged series of edits and subsequent reversals by conflicting parties. The resulting edit pattern is known as an *edit war*. [9] These clashes create a potential for further investigation into the geopolitics of article contribution. Where do the first reports of an event originate? As later iterations of revisions turn these reports into historical accounts, are these editors from the same country? And more generally, to what extent is a collection of these articles written by volunteers located at the respective location of the event?

In this thesis I will propose a method to help answer these questions. By determining the geographic origin of each edit to an article, I will be able to calculate the geographic distribution of contributors. This distribution will then be used to answer the questions above for either a single article or a collection.

Include complete summary and key findings?

1.1 STRUCTURE

³ http://en.wikipedia.org/wiki/Arab_Spring (visited on 01/29/2012)

⁴ http://en.wikipedia.org/wiki/Egyptian_Revolution_of_2011 (visited on 10/31/2011)

⁵ http://en.wikipedia.org/wiki/2011_Libyan_uprising (visited on 10/31/2011)

⁶ http://en.wikipedia.org/wiki/2011_Bahraini_uprising (visited on 01/29/2012)

⁷ http://en.wikipedia.org/wiki/Wikipedia:Neutral_point_of_view (visited on 12/08/2011)

complete over time, name the basic chapters and their function, one part = one paragraph

The chapter **FOUNDATION** provides background information about **Wikipedia**, article editing (**Contributions**), and the application of geographic data (**Georeferences**). The first part ends with **HYPOTHESES** where I propose the research questions that this thesis hopes to answer.

In **APPARATUS** I will describe the tools available to be used in the method that will be applied to a host of articles in **EXPERIMENTS**. The findings will be presented in **RESULTS**. Followed by a discussion of their feasibility in **CONCLUSION**.

FOUNDATION

Wikipedia is a phenomenon that has attracted researchers across all fields, notably computer science and sociology, who have written over 1,000 reports on the subject to date. Nielsen [10] compiled an overview of Wikipedia research¹ and divides these publications into four categories:

CONTENT PRODUCTION Covering all aspects of voluntary production such as motivation, collaboration, coverage and bias, quality and vandalism, actuality, and geography.

INFORMATION USE Treating how the resulting corpus is being used, e.g. Wikipedia citation in research, use in court, trend-spotting, natural language processing and automatic translation tools, thesaurus construction, or categorization.

IMPROVEMENT These are studies concerned with the improvement of both the software used by Wikipedia and the content, e.g. automatic linking, improved editors, as well as quality and trust indicators.

COMMUNICATION Studies in this category look at Wikipedia as an online collaboration tool for education and research.

This thesis falls into the first category, content production, as it examines the geography of article contributions that will become part of the Wikipedia's corpus. After a short overview of Wikipedia from a user's perspective, I will introduce its model of collective authorship and present prior research concerning location and geography.

2.1 WIKIPEDIA

Wikipedia is an online encyclopedia with editions in over 260 languages. Counting 3.6 million articles, the English version is by far the biggest. However, other language editions differ sharply in size and usage.[11] If articles covering the same topic exist in other language editions, they can be connected by inter-language links.²

english wikipedia is focus for this thesis

¹ Another resource is the Wikimedia Foundation's own directory of Wikipedia research projects at <http://meta.wikimedia.org/wiki/Research:Projects> (visited on 10/12/2011).

² As these links are in fact links between different wikis, they are also called *interwiki* links.

2.1.1 *History*

Wikipedia was officially started on 15 January 2001 by Jimmy Wales and Larry Sanger. Wales previously founded Nupedia, a free and peer-reviewed online encyclopedia written only by experts. However, the speed of content production was extremely low. Wikipedia was founded as a feeder project to collectively write on articles before these entered Nupedia's review process. Wikipedia then quickly created other language editions and dwarfed its predecessor³.[\[12\]](#)

After being mentioned on Slashdot, a technology news website, in March 2001, Wikipedia quickly attracted new users. This tech-savvy group of people created new articles at a staggering rate of 1,500 articles per month in the first year. These articles then quickly started showing up in Google's search results, attracting even more new users. The non-English editions grew slower but as a group accounted for 75% of all articles in 2007. By 2011 the combined article count passed 20 million.[\[12\]](#)

2.1.2 *Wikimedia Foundation*

Wikipedia is operated by the Wikimedia Foundation, a non-profit organization founded in Florida on June 20, 2003. It is completely financed by public contributions, such as donations and grants. Individual grants can be quite substantial—among the most generous donors are Google and the Stanton Foundation handing out respectively \$2 million and \$3.6 million in single donations.[\[13\]](#)

In the Wikimedia Foundation's 2010–11 fiscal year, \$8.9 million was spent on website operations, including server hosting and software maintenance. The rest of the \$20.0 million of total expenditures went into complementary activities such as fund raising, administration, and the support of local chapters.[\[14\]](#)

The local chapters are self-dependent organizations set up in countries around the globe to locally promote the foundation's cause and collect donations. The first local chapter to be created was Wikimedia Deutschland, founded in Berlin in 2004.[\[13\]](#)

The individual language editions of Wikipedia are not hosted by the local chapters, however. All of Wikipedia's content is centrally stored on servers in Tampa, Florida and in Amsterdam, Netherlands.[\[13\]](#)

³ Only 24 articles were completed in Nupedia's review process. The project was officially ended in 2003.

2.1.3 Anatomy of an article

All Wikipedia articles share a similar layout: a large content area topped by the article title. Article titles can change over time, e.g. *2011 Bahraini uprising* was renamed to *2011-2012 Bahraini uprising*⁴. For these cases, Wikipedia has a redirecting mechanism that forwards the visitor to the final article and displays a small note below the title (see figure 1).



Figure 1: The article *2011–2012 Bahraini uprising* viewed in a web browser on 01/23/2012.

Occasionally the content section can be topped by one or more warning boxes to inform the visitor that the article is violating an editing policy, e.g. the information of the article may be out-dated because it is subject to current events. When an article spans several sections a table of contents is added below the first introductory paragraphs. In addition to prose, some articles feature info boxes on the right hand side. These boxes show information in a structured way and can be found on articles of similar topics, giving the visitor a quick glance on key information without having to read the text.

This information may include dates and geographic coordinates. E.g. the article *2011–2012 Bahraini uprising* has the time interval “14 February 2011 – ongoing” and is tagged with the coordinates 26°01′39″N 50°33′00″E, pointing to the centre of Bahrain. Even when no coordinates are present in the article, it still may be associated to a location. In that case the info box just presents the place, instead of providing the geographic coordinates, e.g. *Maspero demonstrations*⁵.

⁴ http://en.wikipedia.org/wiki/2011-2012_Bahraini_uprising (visited on 01/23/2012)

⁵ http://en.wikipedia.org/wiki/Maspero_demonstrations (visited on 01/25/2012)

2.1.4 Categories

At the bottom of each article is an optional list of categories that the article belongs to, e.g. the article *2011-2012 Bahraini uprising* belongs, among others, to “Arab Spring by country”⁶ and “2011 protests”⁷.

Categories can not only consist of pages but also of sub-categories, e.g. “Arab Spring by country” has the sub-category “2011 Libyan civil war”⁸ which in turn has 5 sub-categories and 55 pages. When looked at as a graph, the system of categories does not form a tree, however, for there is no restriction on the inclusion of categories—even cycle-creating inclusions are possible. The set of articles in a category can be as arbitrary as its topology. The category “Arab Spring by country” does not only contain articles covering the Arab Spring by country, but also articles about killed individuals, e.g. “Zakariya Rashid Hassan al-Ashiri”⁹.

2.2 MEDIAWIKI AND EDITING

Anyone with a browser and internet access can edit Wikipedia’s articles¹⁰. In collaboration, people all over the world contribute and improve the content. This is made possible by MediaWiki, the server software that makes Wikipedia a wiki. The software allows website visitors to add and modify the page content in the browser using *wikitext*, simplified markup language.¹¹ Its syntax can be used to structure a text into sections, embed images, and links to other pages, similar to HyperText Markup Language (HTML). The syntax is explicitly kept simple to keep the entry barrier to editing low, e.g. adding an article to a category is as easy as putting the category name at the end of the wikitext.

2.2.1 Templates

Wikitext has a special syntax for templates¹². These are reusable containers for text snippets and repetitive material like the info boxes described in [Anatomy of an article](#). When a template

⁶ http://en.wikipedia.org/wiki/Category:Arab_Spring_by_country (visited on 01/23/2012)

⁷ http://en.wikipedia.org/wiki/Category:2011_protests (visited on 01/23/2012)

⁸ http://en.wikipedia.org/wiki/Category:2011_Libyan_civil_war (visited on 01/23/2012)

⁹ http://en.wikipedia.org/wiki/Zakariya_Rashid_Hassan_al-Ashiri (visited on 01/23/2012)

¹⁰ Some articles can be locked because of sustained vandalism or content disputes.^[15]

¹¹ For the syntax see http://en.wikipedia.org/wiki/Help:Wiki_markup (visited on 12/12/2011).

¹² <http://en.wikipedia.org/wiki/Help:Template> (visited on 01/23/2012)

is used in a page, the server software replaces the template placeholder—the template name surrounded by curly brackets—with the template’s content. The content can be parameterized with key-value pairs so that, for example, an info box about countries can be used by several countries’ articles.

By using templates, the information is likely to be more structured than simple, free-form text. Each invocation of a template also renders its content in the same way, allowing for and encouraging more consistency¹³. More importantly, the usage of a template in an article lets that article become a member of the group of articles that embed this template. This is an alternative mechanism to group articles that is likely to yield more homogeneous results than the category system, see [Categories](#).

2.2.2 Revision history

Each submission of an edit in the browser creates a new revision of the article and is stored in the revision history, see figure 2. Naturally, each article available today started from an empty page and is the result of a succession of edits.



Figure 2: Revision history of the article *2011–2012 Bahraini uprising* on 01/23/2012.

Each entry in the revision history consists of the new wikitext, the date of submission, the user name, and an optional comment explaining the change. In addition, authors have to possibility to mark an edit as *minor*. In doing so, they suggest that the submitted change is only superficial, e.g. correcting typography or moving passages of text, and does not change the meaning of the article.^[16]

Each revision can not only be examined by other users but also reverted. Especially in the case of vandalism this mechanism can be used to restore the previous state of the article. Reverts may

¹³ E.g. when an editor notices that an info box supports the parameter “location” but does not have one yet, the user may feel encouraged to complete it.

also appear when different views on the same topic collide. To minimize the potential for *edit wars*[9] Wikipedia urges its users to discuss controversial topics on the article's talk page.

2.2.3 Authors

Contributions to an article can be done anonymously or as a registered user. A registered user gains privileges like the ability to create articles or the use of the social network features in Wikipedia. With the initial registration a *user page* is created where the user is allowed to publish a profile and interact with other registered users.[17] The majority of edits come from registered users; anonymous edits account for a quarter of all edits.[18]

A third group of editors are automatic programs known as *bots*. They perform routine tasks ranging from spell-checking, to curse word detection, to automatic reverts on vandalism. Currently the English Wikipedia alone has nearly approved 1,500 bot tasks running, either automatically or manually triggered by a real user.[19]

2.2.4 User pages

When a Wikipedia user decides to register, a *user page* is created for him on Wikipedia's website. This is a special page that can be edited like any other article. The user can publish personal information in prose, or reuse a template (see [Templates](#)). The templates available to decorate one's user page include the following:

- Spoken languages, e.g. "This user is a native speaker of English."¹⁴
- Location, e.g. "This user comes from India."¹⁵
- Expression of personal views, e.g. "This user opposes Imperialism."¹⁶

Of course, a user can publish just about anything. As a result, the information on a user page may not be accurate.¹⁷

Like any other article, each user page has a discussion page that can be used to communicate with that user by leaving a

¹⁴ <http://en.wikipedia.org/wiki/Wikipedia:Babel> (visited on 01/23/2012)

¹⁵ http://en.wikipedia.org/wiki/Template:User_India (visited on 01/23/2012)

¹⁶ http://en.wikipedia.org/wiki/User:Serouj/UserBox/Against_Imperialism (visited on 01/23/2012)

¹⁷ See user Lihaas, who seems to hail both from India and from Pakistan: <http://en.wikipedia.org/wiki/User:Lihaas> (visited on 01/23/2012)

message. Viégas et al. [20] looked at how contributors coordinate their actions and found that these pages “hold much of the community interaction”.

2.3 CONTRIBUTIONS

Wikipedia’s articles are continuously edited by its users. The nature of an edit can range from simple spelling or grammar correction, to improving the content of a sentence, to writing or removing whole articles. This collective authorship makes it difficult to determine an individual author’s contributions; in other words, it is not easy to tell who wrote what.

Research in this area tends to be motivated by the desire to identify individual authors with a good reputation in order to assign a trust score to them. This is based on the assumption that trusted authors consistently produce high quality contributions that outlive contributions of lower quality. Kramer, Gregorowicz, and Iyer [21] devised a method to assign trust scores to the authors of an article by examining the wealth of information contained in the article’s revision history. They looked at an article as being a set of phrases. The author who first wrote a sentence gets the credit for that phrase and will gain trust if it survives future edits.¹⁸

A similar approach of calculating the longevity of text chunks was followed by Adler and De Alfaro [22]. They adapted standard text-diff algorithms to the peculiarities of the wiki revision system, e.g. keeping track of text chunks that were removed at one point and then reinserted in later revisions. Based on these algorithms a reputation system was implemented by Adler et al. [23] which offers an application programming interface (API)¹⁹ that can annotate a Wikipedia article. The annotated text is the result of splitting the original text into chunks and attributing them with their respective authors, the number of the revision where the chunk was added, and a trust value for the author.²⁰

2.4 GEOREFERENCES

In order to analyze the localness of contributions, it is necessary to geotag them, i.e. applying geospatial metadata like coordinates to each contribution, derived from the author’s location.

¹⁸ Kramer, Gregorowicz, and Iyer [21] define a sentence as an n-gram—a sequence of n words—and use a sliding window model to follow it across revisions to prevent simple rearrangements of text from counting as a new sentence.

¹⁹ <http://www.wikitrust.net/vandalism-api> (visited on 10/31/2011)

²⁰ Adler et al. also released a Firefox add-on that highlights untrustworthy passages when viewing Wikipedia articles: <https://addons.mozilla.org/en-US/firefox/addon/wikitrust/> (visited on 11/15/2011).

In his doctoral thesis Hardy [24] used the Wikipedia corpora to study the spatial behavior of article production. The dataset was limited to anonymous users and articles that were geotagged.

For each anonymous contribution an Internet Protocol (IP) address, belonging to the point of Internet access, is stored in the revision that is created. Various methods to determine the geographic location from a given IP address have been studied by Muir and Oorschot [25]. Various visualizations^{21 22} of edit distributions use geolocation databases like MaxMind²³ and Quova²⁴.

For registered users, the IP address is not stored with the revision. Therefore IP geolocation services cannot be used. Lieberman and Lin [26] found a novel approach by assuming users prefer to edit geographic articles in their proximity. The approximated user location was derived from the center of the convex hull around those articles.

The user pages offer another source for finding the author's location. Entity names like a city or a country can be extracted from both the prose or the info boxes that the users put on their page.

Building on this foundation, I will formulate a set of **HYPOTHESES** in the next section.

²¹ <http://infodisiac.com/blog/2011/05/wikipedia-edits-visualized/> (visited on 10/31/2011)

²² <http://sonetlab.fbk.eu/wikitrip/> (visited on 10/31/2011)

²³ <http://www.maxmind.com> (visited on 10/31/2011)

²⁴ <http://www.quova.com> (visited on 10/31/2011)

HYPOTHESES

To gain insight on how Wikipedia is being used during and after political events¹, and ultimately whether articles covering the events are written by people that are most affected, I will propose a set of hypotheses aimed at different aspects of article production.

For this thesis I will use the event definition proposed by Lewis [27, p. 243]:

“An event is a localized matter of contingent fact.
[...] An event occurs in a particular spatiotemporal region.”

It follows that events must have clear spatial and temporal boundaries. The spatial boundaries give it a location, distinguishing the where. The temporal boundaries, namely the start and the end date, divide the event into three intervals: *before*, *during* and *after*.² Let me further denote the *beginning* of an event as the first seven days of an event.³ In conjunction with location this division into time intervals allows for a detailed look into [Article creation](#), the level of [Participation](#) in subsequent intervals as well as [Text survival](#).

3.1 ARTICLE CREATION

The use of Wikipedia during a political event starts with the creation of the article describing the event. How quickly a users create such articles is less clear. Wikipedia’s popularity and global audience provide an incentive to publish without delay. This leads to the following hypothesis:

HYPOTHESIS 1 (H1). Articles are created with only a short delay after the start date of the event.

Moreover, considering the rising year-on-year Wikipedia usage in numbers of pages viewed[6], this delay should become shorter and shorter suggesting an increased use of Wikipedia as a news channel, thus:

- ¹ I will not try to argue what makes an event political but rather identify a set of events by picking a suitable category of articles.
- ² For ongoing events the end date will be the date of the analysis.
- ³ The interval was picked arbitrarily but acknowledges the fact that event dates in Wikipedia articles rarely carry a time attribute, therefor a shorter interval, say 24 hours, is less feasible.

HYPOTHESIS 2 (H2). The more recent an article, the shorter is the delay between the event start and article creation.

A user has the chance to create a new article in any of the 260-odd language editions of Wikipedia. Although the English version is by far the biggest and most visited, it would be interesting to see whether it is also the first choice to create an article covering a new event. Shortly after the first article has been created, articles covering the same topic are produced across various editions of Wikipedia. These articles are then being linked, mostly manually, via interlanguage links.⁴ When studying knowledge diversity across language editions, Hecht and Gergle [29] found that the English edition is not the superset of concepts of all editions as was previously believed. This means authors retain knowledge they consider important only to their compatriots. For a citizen of a country where English is not an official language, creating an article becomes a political decision: should the author make the information available to fellow citizens or to a world-wide readership? The following hypothesis will be tested only for articles where the user faces that choice, i.e. for articles about events in countries where English is not an official language, thus:

HYPOTHESIS 3 (H3). Articles are being created first in the English Wikipedia.

Regarding the localness of contributions, Hardy [24, p. 57] has established that Wikipedians write about places in their proximity more often than distant ones. His sample included only articles that have a geotag. Naturally, articles about geographic places like towns and sites⁵ will dominate this sample. Since my thesis is concerned with political events, I find this point to be worth revisiting⁶, as the people most affected should be the ones creating the article, thus:

HYPOTHESIS 4 (H4). Articles about political events are created by people in the events' proximity⁷.

⁴ Adar, Skinner, and Weld [28] found that between two languages the interlinking is not symmetrical, i.e. the number of out-links does not match the number of in-links. Links are either missing on one side or the respective topics are not congruous and the user intentionally left out one direction.

⁵ According to Kittur, Chi, and Suh [30] articles about "geography and places" are third biggest group.

⁶ In addition, Hardy [24, p. 61] considered only anonymous users. Since creating an article is only allowed for registered users, his method has to be extended.

⁷ Hardy [24, p. 57] defined proximity not in absolute terms, rather he considered the likeliness of authors being located less far than the average distance between an article and all its contributors.

Hypotheses 1–4 will only be tested against articles that were created as a reaction to an event that has already started. This excludes scheduled events like elections, e.g. *Russian legislative election, 2011*⁸ which was created 335 days before the election date, almost a year in advance.

3.2 PARTICIPATION

A Wikipedia article usually has more than one author. Once it has been created, users from around the globe can edit an article collectively. Viégas, Wattenberg, and Dave [31] tried to find patterns in the revision history that would reveal certain aspects of collaboration or the lack thereof, e.g. discussions and vandalism.⁹ In respect to authorship, the researchers found the proportions of anonymous contributions differed strongly from page to page while showing no preference to any topic. This inconclusive result and the age of the sample¹⁰ merits further investigation.

In 2007, Kittur et al. [32] found that a core of registered users is still doing the bulk of all edits. However, anonymous users contribute considerable amounts of text. For accounts of political events, due to their dynamic nature, I expect a strong participation by unregistered users while the events are still unfolding:

HYPOTHESIS 5 (H5). In the beginning of the event anonymous users contribute more than registered users.

HYPOTHESIS 6 (H6). For the duration of the event there are more local contributions than distant ones.

When the political event is considered “over” its end date in the article changes from “present” to a calendar date. In this unbounded and final phase I would expect the flood of contributions to subside and the content to consolidate when editors tighten prose or remove text they believe to be irrelevant. Looking at the whole lifespan of an article I would also expect registered users to outnumber anonymous ones as suggested by Kittur et al. [32] and the spatial distribution of contributions to become less local. Thus the final hypotheses:

HYPOTHESIS 7 (H7). Articles of political events that have ended tend to shrink in size.

⁸ http://en.wikipedia.org/wiki/Russian_legislative_election,_2011 (visited on 01/07/2012)

⁹ Using their history flow visualization Viégas, Wattenberg, and Dave [31] first identified patterns in single articles and later tried to statistically confirm their prevalence by analyzing the complete English corpus.

¹⁰ Viégas, Wattenberg, and Dave [31] used a dataset from May 2003.

HYPOTHESIS 8 (H8). After an event has ended, there will be more contributions from registered users than from anonymous ones.

HYPOTHESIS 9 (H9). After an event has ended, the spatial distribution of the contributors will become less local.

3.3 TEXT SURVIVAL

In 3.2 the contributions are only treated in volume, giving credit to each contributor. However, when multiple authors write the same article, they do not only add text but also modify or even delete parts. A user who reads an article will only see the text that has survived all edits after it was added. Viégas, Wattenberg, and Dave [31] found that early contributions have a high survival rate. Recognizing this *first-mover advantage*, I suspect that accounts of political events show a strong localness in the beginning. Thus the key hypotheses from 3.2 have to be extended to reflect the spatial distributions of the contributions that make up the article:

HYPOTHESIS 10 (H10). For the duration of the event the article text contains more local contributions than distant ones.

HYPOTHESIS 11 (H11). After an event has ended, the spatial distribution of the surviving contributions will become less local.

This concludes the statement of the hypotheses. The next chapter, APPARATUS, describes an application designed to test these.

Part II

METHODS

4

APPARATUS

This chapter describes the application designed to test the hypotheses. The first section describes the [Data sources](#) and how they can be used to provide the content for the article analysis. In [Application design](#), I will give an overview of the application's architecture as well as the technologies being employed. The section [Algorithms](#) describes with how an article's eligibility for analysis is being determined as well as other key methods regarding information extraction, e.g. date and location parsing. In [Visualization](#) I give a quick introduction into the charts being used to present the data, and that will be featured more prominently in the following chapter [EXPERIMENTS](#).

4.1 DATA SOURCES

For an automated analysis, simply browsing Wikipedia's website is not really feasible. The bulk of Wikipedia's content, e.g. articles, revisions, and discussions, is stored on its database servers. Unfortunately, these databases are not directly accessible over the Internet. The Wikimedia Foundation, however, makes a lot of the data available in the form of database dumps or through an API.

4.1.1 *Wikipedia website*

For a complete article analysis, navigating the website can be tedious, as one would have to click through a complete revision history and parse the page's source which is formatted in HTML. However, individual pages contain data that is static and can be used throughout the analysis process. This makes it worth writing a specific parser for a technique known as *screen scraping* to extract the information. On a high level, it involves the following steps:

1. Looking at the HTML source of the page and identifying how the HTML tags and attributes that are used to structure the information.
2. Writing a parser that addresses the identifying tags and thereby tokenizes the data.
3. Converting the found tokens into an output format, e.g. JavaScript Object Notation (JSON).

For a simple HTML table, a parser can be written in a few lines of code. Using this technique, the following static information was gathered:

BOTS A list of bots was built based on the Wikipedia page *List of bots by number of edits*¹. This list is used to distinguish bots from real authors, as contributions done by bots are excluded from the analysis. There are unregistered bots, however, which do not appear on the list. For a lack of automated distinction, these are counted as normal authors.²

COUNTRIES A list of countries was extracted from the article *ISO_3166-1*³. It provides a list of standardized country names that is also respected by Wikipedia's authors when referring to a country by name. In a second pass, the Wikipedia article of each country was retrieved to extract the country's for geographic coordinates⁴. For a discussion on countries and coordinates, see [Geographic resolution](#).

OFFICIAL LANGUAGES For each entry in the list of countries, the country's official languages were added. From the *List of ISO 639-1 codes*⁵, all languages including their two-letter codes were extracted. In a second step, the *List of official languages by state*⁶ was scraped and merged with first list by matching the language name.⁷ The result is a mapping from a country's name to a list of two-letter codes of that country's official languages.

Making both of these sets static was a design decision recognizing the trade-off between having them in memory and querying for each article.

¹ http://en.wikipedia.org/wiki/Wikipedia:List_of_bots_by_number_of_edits (visited on 01/24/2012)

² A simple heuristic employed by other software to analyze MediaWiki content is treating all contributors whose username contains or whose comments start with "bot" as a bot, e.g. pymwdat, see <http://code.google.com/p/pymwdat/source/browse/trunk/toolkit.py?spec=svn13&r=13> (visited on 01/24/2012). This has a potential for false positives and is not used.

³ http://en.wikipedia.org/wiki/ISO_3166-1 (visited on 01/02/2012)

⁴ For some countries, coordinates were not present on the page, e.g. <http://en.wikipedia.org/wiki/Australia> (visited on 01/24/2012). In that case, they were manually added by using that country capital's coordinates.

⁵ http://en.wikipedia.org/wiki/List_of_ISO_639-1_codes (visited on 01/29/2012)

⁶ http://en.wikipedia.org/wiki/List_of_official_languages_by_state (visited on 01/29/2012)

⁷ Not all countries have official languages—including the United States of America. In that case, the *de facto* official languages were used, e.g. English and Spanish for the United States of America.

4.1.2 Database dumps

Monthly database snapshots of all wikis run by the Wikimedia Foundation, including Wikipedia, are publicly available⁸ as database dump files in the Extensible Markup Language (XML) file format. For each of the wikis a variety of dumps is available that include all articles and, optionally, their revision history, all categories, interlanguage links, etc. Despite this openness, some database tables are not publicly available. The dump files of the database tables *users* and the *watchlist* are kept private.

The dump files can be quite large, e.g. a compressed dump of all articles of the English Wikipedia in their current revision has a size 7.3 GB.⁹ This huge size makes processing them rather slow.¹⁰ When analyzing only a single article or a category of articles, the MediaWiki API can deliver the same information contained in the dumps in a much more targeted manner.

4.1.3 MediaWiki API

Wikipedia runs on the open source software MediaWiki, written in PHP: Hypertext Preprocessor (PHP). It offers a well documented API¹¹ which can be used by other programs to remotely use the wiki's features such as changing content and restoring revisions.¹² For analysis of articles, the API offers queries directed at a variety of article properties, e.g. revisions, categories and, links. Among the output formats for the responses are JSON and XML. Similar to MediaWiki's Special:Export page¹³, the API also offers an article export that includes all revisions.

Listing 1: Example JSON response to a query to list all bots that edited the article *2011-2012 Bahraini uprising*

```
{
  "query": {
    "redirects": [{
      "from": "2011 Bahraini uprising",
      "to": "2011-2012 Bahraini uprising"
    }],
    "pages": {
```

⁸ <http://dumps.wikimedia.org> (visited on 12/11/2011)

⁹ The uncompressed size is 31.0 GB, see http://en.wikipedia.org/wiki/Wikipedia:Database_download (visited on 12/11/2011)

¹⁰ The project WikiHadoop addresses this problem by offering a stream task format to be used in Hadoop (MapReduce) infrastructure, see <https://github.com/whym/wikihadoop> (visited on 12/11/2011).

¹¹ http://www.mediawiki.org/wiki/API:Main_page (visited on 01/24/2012)

¹² The full capability of the API can be seen at tried at the *Sandbox* at <https://en.wikipedia.org/wiki/Special:ApiSandbox> (visited on 01/24/2012), a recent addition to the MediaWiki software.

¹³ The page <https://en.wikipedia.org/wiki/Special:Export> (visited on 12/11/2011) allows for exporting of articles from the English Wikipedia.

```

    "30876395": {
      "pageid": 30876395,
      "ns": 0,
      "title": "2011-2012 Bahraini uprising"
    },
    "allusers": [{
      "userid": "13146235", "name": "28bot"
    }, {
      "userid": "5415725", "name": "718 Bot"
    }, ..., {
      "userid": "13770078", "name": "AWBCPBot"
    }
  ],
  "query-continue": {
    "allusers": {
      "aufrom": "AWeenieBot"
    }
  }
}

```

Some of the queries have a limit on how many results they return on a single request. When there are more results, the response contains a *query-continue* attribute that can be sent with following query so that the next result set can be returned. The following API calls will be important for this thesis:

QUERY INFO This basic query is returns essential information like the article ID, the last revision ID, but also the full wiki-text of the last revision.

QUERY REVISIONS Lists all revisions for an article and for each includes a timestamp and the user as well as the comment for the text change.

QUERY CATEGORYMEMBERS For a given category, this query lists the articles and subcategories belong to it. This query can be used to construct groups of articles for analysis in this thesis.

QUERY TEMPLATEEMBEDDERS For a given template name, this request lists all pages that embed it. This query can also be used to build a group of articles.

OPEN SEARCH A method to suggest articles, categories, and templates that contain a term. It can be attached to an input field where a user is supposed to enter the name of an article.

PARSE This special query returns the HTML version of the article's wikitext. The content that is returned is exactly the HTML source that is sent to a browser when a visitor looks at this article or user page. This query will be used in cases where it is easier to parse the HTML markup than the wikitext, e.g. the [User pages](#).

4.1.4 Toolserver

The Germany based Wikimedia Deutschland e.V. runs Toolserver¹⁴, a platform for software tools that can access a continuously updated copy of Wikipedia's databases. Among these replicated databases is the English Wikipedia and other major language editions. However, the deployment of self-made software scripts is restricted and requires an account on Wikimedia's Toolserver.¹⁵

Some scripts that are already deployed can be accessed freely, allowing them to be reused. One of these was developed by SoNet¹⁶, a social networking research group based in Italy, for a project called WikiTrip (see [Analysis projects](#)). It offers an API¹⁷ to get simple article statistics like article ID, text length, as well as complex data structures like a list of unique editors including their gender, if they are registered users and chose to reveal their gender in their Wikipedia account.¹⁸

In effect, calling the SoNet API replaces several calls to the original MediaWiki API and thereby speeds up the information retrieval, especially when the number of revisions or authors is high. The returned data object has the following structure:

Listing 2: SoNet API response to a query for the article *2011-2012 Bahraini uprising*

```
{
  "first_edit": {"timestamp":1297734917,"user":"Master&Expert"},
  "count":1778,
  "minor_count":401,
  "count_history":{"today":3,"week":5,"month":90,"year":1778},
  "last_edit":1327370324,
  "totaldays":0,
  "average_days_per_edit":"0.00",
  "edits_per_month":0,
  "edits_per_year":0,
  "edits_per_editor":"4.17",
  "editor_count":426,
  "anon_count":337,
  "editors": {"Bahraini Activist":
    {"all":106,"minor":21,"first":"17 May 2011, 09:45:25",
    "last":"22 January 2012, 10:52:50","atbe" 203811,
    "minorpct":"19.81", "size":"140.54","urlencoded":"
      Bahraini_Activist"},
    ...
  },
  "anons":{"2011-02-15T08:11:52Z":
```

¹⁴ <http://toolserver.org> (visited on 12/11/2011)

¹⁵ I applied for a Toolserver account outlining my necessary database queries, usage profile, as well as my affiliation with the Freie Universität Berlin. The application was submitted on 2011-12-21 and has not been processed yet (2012-01-23).

¹⁶ <http://sonetlab.fbk.eu/> (visited on 12/12/2011)

¹⁷ The API is documented here: <https://github.com/volpino/toolserver-scripts/tree/master/\ac{PHP}> (visited on 12/12/2011)

¹⁸ Try http://toolserver.org/~sonet/api_gender.\ac{PHP}?article=Egypt&lang=en (visited on 12/11/2011) to get a list of all registered users who edited the article *Egypt* of the English Wikipedia.

```

        ["78.2.29.139", "Rovinj Croatia", 45.08, 13.64],
        ...
    }
}

```

This high density of preprocessed information shows the power of the Toolserver and its direct access to the database. The property *editors* lists all unique authors of an article and their edit count (property *all*). The second exhaustive collection is under the property *anons*. There, all anonymous contributors are listed with their IP address, as well as their geographic region and coordinates. The geographic lookup is uses¹⁹ the GeoCityLite database from Maxmind (see [IP Look-up](#) for a discussion).

4.1.5 Third-party sources/Web services

Like the Toolserver scripts in the previous section, other research projects exist that can be reused as data sources. Depending on the project's goal, a variety of preprocessed data is available:

ARTICLE TRAFFIC Wikipedia user Henrik²⁰ provides a web service that processes Wikipedia's log files²¹ to calculate the number page views per article for a given time. These statistics can be viewed through a browser²² or queried through an API²³.

CATSCAN This web service, offered by Toolserver administrator Duesentrieb²⁴, finds articles that belong to a given category and its sub-categories (see [Categories](#) on why this is non-trivial). It also offers to limit the search to an intersection of categories, e.g. German politicians who are also physicists²⁵. The results are presented in the browser or can be downloaded as a file in the comma-separated values (CSV) format.

¹⁹ <https://github.com/volpino/toolserver-scripts/blob/master/\ac{PHP}/api.\ac{PHP}> (visited on 01/24/2012)

²⁰ <http://en.wikipedia.org/wiki/User:Henrik> (visited on 12/12/2011)

²¹ These are available at <http://dumps.wikimedia.org/other/pagecounts-raw/> (visited on 12/12/2011)

²² E.g. http://stats.grok.se/en/201105/2011_Egyptian_Revolution (visited on 12/12/2011)

²³ E.g. http://stats.grok.se/json/en/201105/2011_Egyptian_Revolution (visited on 12/12/2011)

²⁴ <http://meta.wikimedia.org/wiki/User:Duesentrieb> (visited on 12/13/2011)

²⁵ [https://toolserver.org/~daniel/WikiSense/CategoryIntersect.\ac{PHP}?wikilang=de&wikifam=.wikipedia.org&basecat=Politiker+\(Deutschland\)&basedeep=3&mode=cs&tagcat=Physiker&tagdeep=2](https://toolserver.org/~daniel/WikiSense/CategoryIntersect.\ac{PHP}?wikilang=de&wikifam=.wikipedia.org&basecat=Politiker+(Deutschland)&basedeep=3&mode=cs&tagcat=Physiker&tagdeep=2) (visited on 12/13/2011)

POOR MAN'S CHECKUSER The project *Poor Man's Check User*²⁶ mapped registered users to IP addresses based on a bug in the session management of the MediaWiki software.²⁷ For the period the bug has been active, some usernames could be mapped. Naturally, the more edits a user did in this period, the more likely is an appearance in this list. For the purpose of this thesis, I screen-scraped the entire table and condensed²⁸ it to 14,171 unique users.

QUOVA This geo-location web service maps an IP address to a geographic location, see [Georeferences](#).

WIKITRUST Based on Adler et al. [23] an open source online reputation system²⁹ was set up by the University of California, Santa Cruz, to allow for easy vandalism detection (see [Contributions](#)). Given an article ID and a revision ID, the API method *wikimarkup* returns an annotated version of the wikitext of that revision. An annotation consists of a trust value, the revision ID the text got introduced into the article as well as the authors user name or IP address, e.g. revision 473029564 of the article *2011-2012 Bahraini uprising*³⁰:

Listing 3: Excerpt of the annotated markup for the revision 473029564 of the article *2011-2012 Bahraini uprising*

```
{{#t:7,468889105,Kudzu1}}The
{{#t:7,470041169,Happysailor}}2011-2012
{{#t:8,413989516,Master&Expert}}Bahraini
{{#t:8,427545590,Kudzu1}}uprising, sometimes called the
{{#t:9,455029613,Sitrawi86}}February 14 Revolution
```

All wikitext following an annotation, up to the next one, was written by that author. The web service provider implemented a custom diff algorithm for the attribution of

²⁶ Project website: <http://wikiwatcher.virgil.gr/pmcu> (visited on 01/02/2012). The project's name is a reference to the *checkuser* permission that a community-elected group of registered users possesses. It allows a de-masking of the IP addresses for each of a registered user's edit.

²⁷ When a user exceeded a certain time while editing an article without submitting the current changes, the user's session expired on the server. When the edit was submitted after the expiration the user appeared as an anonymous author, being only known by his IP address. When the user then logged in again, the same change was sent again. Scanning all revisions for the same change set therefor allowed for a matching between user name and IP address. This loophole has been closed, however.

²⁸ Some usernames have multiple entries as each occurrence of the bug created a unique "evidence". Among those, some have been manually verified and ranked. When multiple entries exist, my algorithm picks the top ranked.

²⁹ <http://www.wikitrust.net/vandalism-api> (visited on 10/31/2011)

³⁰ <http://en.collaborativetrust.com/WikiTrust/RemoteAPI?method=wikimarkup&pageid=30876395&revid=473029564> (visited on 01/23/2012)

authorship. This was needed to overcome wiki-specific issues and to maximize tracking, e.g. for text that is removed and re-inserted at a later revision.³¹

4.2 AVAILABLE TOOLS

To process the data from all the data sources, a wide range of software tools are available in the open source community. A simple search for “Wikipedia” on GitHub³², a source code exchange platform, shows a multitude of small software projects. These come in different programming languages and different feature sets and usually help in downloading articles in batches and extract data from big dump files. Developed by vigilantes and researchers alike, these programs facilitate both data retrieval and processing.

4.2.1 *Toolkits*

A group of openly available software packages³³ qualify as swiss-army knives for processing and analyzing [Database dumps](#):

PYWIKIPEDIA As the mother of all Python toolkits, the Python wikipedia robot framework³⁴ offers an extendable set of classes for all MediaWiki entities like a page, a user, revision, etc, and is typically used to write a bot program for automated editing tasks (see [Authors](#)).

PYMWDAT Based on PYWIKIPEDIA, this toolkit offers a convenient downloader for all revisions of an article as well as an extensible dump file analyzer with support for filtering and revert detection.

LEVITATION A creative project to turn [Database dumps](#) into a Git³⁵ repository. As a source code management system, Git offers a more space efficient way to store the sequence of revisions of the articles, since it only stores the difference between revisions. Once converted to a repository, moving from revision to revision is much faster than processing the large dump files. Git’s diff mechanism together with

³¹ <http://www.wikitrust.net/frequently-asked-questions-faq#TOC-On-text-author-and-origin> (visited on 01/24/2012)

³² <https://github.com/search?q=wikipedia&type=Repositories> (visited on 12/12/2011)

³³ Although none of these were used for the content analysis of this thesis, their study proved very insightful on how to process Wikipedia’s content.

³⁴ <http://pywikibot.sourceforge.net/> (visited on 01/24/2012)

³⁵ <http://git-scm.com/> (visited on 01/24/2012)

its *blame* command can be used as an alternative way to attribute authorship to passages of article content.³⁶

4.2.2 Analysis projects

In addition to the toolkits, a handful of research projects exist that process Wikipedia's content. These are purpose built applications that have a much narrower focus but are very skillful in combining and using different data sources such as [MediaWiki API](#), the [Toolserver](#) or other [Third-party sources/Web services](#):

WIKIPRIDE Python web-application³⁷ to visualize contributions of groups of editors that registered in the same month.³⁸

WIKI TRIP JavaScript application³⁹, written at SoNet, that uses the [MediaWiki API](#) as well as its own [Toolserver](#) scripts, to visualize the evolution of a single article over time including: anonymous vs. registered contributors, male vs. female registered users, anonymous edits by country.⁴⁰

4.3 APPLICATION DESIGN

For the analysis of articles I developed an application that could draw data from the different sources and process that data in a timely fashion. Following the impressive WikiTrip application design (see [Analysis projects](#)) I decided to build a web application that runs entirely in a web browser.⁴¹

4.3.1 Technologies

As a web application the software heavily relies on HTML version 5, JavaScript (JS), and Cascading Style Sheets (CSS). It uses a range of open-source toolkits and libraries for a variety of purposes:

BOOTSTRAP Twitter's web application toolkit⁴² controls the basic layout, the styling of sections and form fields, and the navigation bar at the top.

³⁶ In fact, git operates on a line level, making the attribution rather coarse. To get blaming functionality on a word level, I patched the source, see my fork at <https://github.com/davkal/levitation/commit/5fca0001d26cb67fde6ff9d8a5f2b1414cf7681e> (visited on 01/24/2012).

³⁷ <https://github.com/declerambaul/WikiPride> (visited on 12/11/2011)

³⁸ Project website: <http://meta.wikimedia.org/wiki/Research:WikiPride> (visited on 12/11/2011)

³⁹ <https://github.com/volpino/wikipedia-timeline> (visited on 12/11/2011)

⁴⁰ Live demo: <http://sonetlab.fbk.eu/wikitrip/> (visited on 12/11/2011)

⁴¹ Application development was done mainly using the Google Chrome browser. The application should run in any HTML5-capable browser.

⁴² <http://twitter.github.com/bootstrap/> (visited on 01/25/2012)

JQUERY The JS library jQuery⁴³ is used to asynchronously retrieve data from the various data sources, dynamically insert elements into the layout as well as parsing screen-scraped webpages by making use of its selectors for addressing elements. The AUTOCOMPLETE⁴⁴ widget of jQuery's UI library is used to display article suggestions based on what has been entered into the article search field.

UNDERSCORE Underscore⁴⁵ is a “utility-belt library” for functional programming in JS. Its *map* and *groupBy* methods are being heavily used in the analysis of content.

BACKBONE Built upon Underscore, Backbone⁴⁶ provides a way to structure a JS application. All models, collections, and views in the application are encapsulated by Backbone objects that can communicate with each other via events.

DATEJS Date.js⁴⁷ is a library that was being used to parse dates in articles about events.

LZ77 A JS implementation⁴⁸ of the LZ77 text compression algorithm⁴⁹. It is used to shrink the size of results when stored in the browser's limited data store⁵⁰.

D3.JS Data-driven documents (d3)⁵¹ is a library to visualize big data sets. This application uses d3's box plots⁵² to summarize the quantitative distribution of analysis results.

GOOGLE CHART TOOLS Google's⁵³ chart tools⁵⁴ provide a wide range of chart types that I used in the application, including line chart, scatter chart, and motion chart.

4.3.2 Models

Looking at Wikipedia's website as a system, visible items like pages, authors, and revisions can be abstracted into classes of

43 <http://jquery.com/> (visited on 01/25/2012)

44 <http://jqueryui.com/demos/autocomplete/> (visited on 01/25/2012)

45 <http://documentcloud.github.com/underscore/> (visited on 01/25/2012)

46 <http://documentcloud.github.com/backbone/> (visited on 01/25/2012)

47 <http://www.datejs.com/> (visited on 01/25/2012)

48 <https://github.com/olle/lz77-kit> (visited on 01/25/2012)

49 http://en.wikipedia.org/wiki/LZ77_and_LZ78 (visited on 01/25/2012)

50 See <http://dev.w3.org/html5/webstorage/> (visited on 01/25/2012) for the limitations.

51 <http://mbostock.github.com/d3/> (visited on 01/25/2012)

52 <http://mbostock.github.com/d3/ex/box.html> (visited on 01/25/2012)

53 Although not open source, some of the charts are developed by the open-source community.

54 <http://code.google.com/apis/chart/interactive/docs/index.html> (visited on 01/25/2012)

objects that can be modeled using BACKBONE (for a complete overview of all models, see *model chart in appendix*).

When a model is instantiated, it knows where to retrieve the data that will populate its attributes, e.g. a revision collection knows that it can get all revisions of an article's history using the query *revisions* from the [MediaWiki API](#). Once the data is retrieved, the attributes of each individual revision object are set, which in turn triggers an event telling a single revision model fetch the annotated wikitext from the WikiTrust API (see [Third-party sources/Web services](#)). While some events trigger the retrieval of more detailed information, others indicate that a model's data is fully populated and ready for display in the application.

4.3.3 Views

The rendering of models is encapsulated in views. Following the publish-subscribe pattern, a view listens to changes in a certain model. When a change event is observed, the view renders itself. For example, the *article view* is rendered multiple times because it draws data from different sources (the page ID is available sooner than, say, the first sentence of the article). Most views, however, rely on several models as they analyze various aspects of an article and then display the results in a chart (see [Visualization](#)).

4.3.4 Article analyzer

This section describes an application run on a high level. For noteworthy algorithms in the sub-routines, see [Algorithms](#).

When loaded in a browser, the application expects an article title, a category or a template name as input. In the case of an article being entered, a click on "Analyze!" starts the following routine, querying the various [Data sources](#):

1. QUERY INFO is being called to check if the article title is a valid title. A successful query also returns the wikitext of the article's latest revision.
2. A call to PARSE retrieves the HTML version of the latest revision as well as the article's links to other language editions. Both the wikitext and the HTML are then being parsed for location and dates.
3. The SoNet Toolserver script is called to retrieve all authors of the article.
4. For all registered usernames in the author collection, a sub-routine is started to locate the users.

5. QUERY REVISIONS is being called to retrieve all revisions (excluding wikitext).
6. The a subset of the revisions, the annotated wikitext is loaded using WikiTrust API.
7. For each day of the article's existence, the page view statistics are loaded.
8. For each language present in the language link collection, the first revision is loaded.
9. All loaded data is analyzed and the results are stored and rendered in the browser.

Depending on the article's number of revisions, this process takes around one minute and involves 100–500 API calls, most of which are done in parallel. During this process, data is rendered whenever it has arrived and been processed.

4.3.5 *Group analyzer*

When a category or template name has been entered, the application switches to *group mode*. On a high level, the group analysis works as follows:

1. Fetch a list of articles by calling QUERY CATEGORYMEMBERS or QUERY TEMPLATEEMBEDDERS.
2. Run the article analyzer routine on all articles.
3. Compute group results and show them in the browser.

During the group analysis only the article overview will be rendered while all of the article analyzer's steps are being executed. While QUERY TEMPLATEEMBEDDERS returns a one-dimensional list, QUERY CATEGORYMEMBERS can return sub-categories. For a category that has been entered into the input field, the pages of its first-level sub-categories will be retrieved as well.

4.4 ALGORITHMS

The individual algorithms described mostly deal with extracting data from the various [Data sources](#). To classify an article as treating an event, the article had to be parsed for a location and a date or date interval. All algorithms will be written in a Python-esque pseudo-code.

4.4.1 *Article requirements*

The presence of a date and a location is a direct requirement for an article to qualify for further analysis when the application is in group mode. All of the following criteria have to be satisfied:

- The article has a location in the form of geographic coordinates or by name, e.g. Cairo.
- The article has a date, e.g. 9 November 1989, or a date interval, e.g. May 2007 - August 2007.
- The article was created after the event started (with a 3 day tolerance) to sort out events that have been scheduled.
- The event (start) date is not before 2002, to make sure that Wikipedia was available as a medium.
- The article does not use certain templates or is part of certain categories, e.g. *Category:Living People*. This has been included for the purpose of filtering out articles that passed the previous tests but are clearly not treating an event.
- The article has at least 100 revisions.
- The article has at least 100 unique authors.
- At least 25% of authors can be located.

Even when all of the requirements above are met, some articles may not have a full set of results after their analysis. They will however be included in the computations for which results are present, see [Hypothesis analysis](#).

4.4.2 *Date parsing*

The way dates are mentioned in articles are as diverse as the people that write them. And when it comes to date intervals, e.g. May 8–12 2007, even a specialized date parsing library like DATE.JS can only be of limited use. I wrote a custom parser that for a given text returns first occurrence of an interval or a single date, i.e. if no interval could be found, the parsing is repeated for a single date.

Some articles embed info box templates (see [Templates](#)) that produce annotated markup using the HCARD⁵⁵ microformat. The

⁵⁵ <http://en.wikipedia.org/wiki/HCard> (visited on 01/25/2012)

annotations can be easily addressed with CSS selectors⁵⁶ and the values conform to the standard for the representation of date and time, International Organization for Standardization (ISO) 8601⁵⁷. In listing 4 the first block tries to parse the microformat annotations. When they are not present, the first info box is checked for a date field. As a last resort, the first sentence and paragraph are scanned for dates to not filter out eligible articles like *2007 Georgian demonstrations*⁵⁸

Listing 4: Date candidates algorithm

```

1 if 'dtstart' in HTML: # checking for hcard
2     start = datejs.parse(dtstart.text)
3     if 'dtend' in HTML:
4         end = datejs.parse(dtend.text) # proper interval
5     elif 'ongoing' in dtstart.next.text:
6         end = today() # ongoing event
7     else:
8         end = start + 1 # single day event
9 elif 'date' in templates.infoboxes[0]:
10     start, end = custom.parse(templates.infoboxes[0].date) # wikitext
11 if not start:
12     start, end = custom.parse(article.first_sentence) # HTML
13 if not start:
14     start, end = custom.parse(article.first_paragraph) #HTML

```

The custom parser then checks for a range of formats including lazy ones like December 14–19, 2008. Over a long iterative process, I identified the following tokens from which to construct the date patterns as regular expressions:

Listing 5: Date tokens

```

1 ords = ['th', 'st', 'nd', 'rd'];
2 tokens = {
3     M: "{0}".format(MonthNames.join('|')), # months
4     D: "\\d{1,2}" + "{0}?".format(ords.join('|')), # day
5     Y: "\\d{4}", #year
6     T: "(/|-|--|sto\\s|\\sand\\s)", # interval delimiter
7     O: "('ongoing'?'*present'*)", # ongoing event
8     F: "From", # ongoing event
9     S: "[\\s]*", # whitespace
10    P: "\\|", # pipe

```

⁵⁶ This requires the template to be properly used by the authors, e.g. [http://en.wikipedia.org/wiki/Eastern_Front_\(World_War_II\)](http://en.wikipedia.org/wiki/Eastern_Front_(World_War_II)) (visited on 01/29/2012) where the template *Start date* is used for the start and the end date, producing the two dates with the same annotation: *dtstart*. Another source of error is the order of year, month, and date, e.g. “Municipal Library Elevator Coup” happened on 28 January 1908 which was added as *Start date|1908|28|01* to the info box but unpredictably rendered as April 1, 1908, see http://en.wikipedia.org/wiki/Municipal_Library_Elevator_Coup (visited on 01/25/2012).

⁵⁷ http://en.wikipedia.org/wiki/ISO_8601 (visited on 01/25/2012)

⁵⁸ The article does not contain an info box and a date is only mentioned in the second sentence: “The demonstrations peaked on November 2, 2007,...” with more dates to follow in the same paragraph. This particular example already shows how idiosyncratic dates can be codified. See http://en.wikipedia.org/wiki/2007_Georgian_demonstrations (visited on 01/25/2012).

```

11 A: "([^- -]*)"# other text
12 };

```

Using these tokens, I produced patterns to match all encountered date formats, e.g. the pattern for December 14–19, 2008 is MDTDY, or to capture the even less conform “From 15 October 2011” (meaning the event is ongoing) the pattern is FDMY. The patterns are ranked by accuracy so that “12 May 2001 - present” is matched before “May 2001 - present”. The coarsest pattern to match is Y, a single year (4 digits).

When a date was extracted using my custom parser, its accuracy is also stored as the *date resolution* with one of the following values: *day*, *month*, *year*.

4.4.3 Location parsing

Like dates, locations and even coordinates can be codified in numerous ways. Most of the coordinate templates used in the info boxes produce annotated markup, thereby making the coordinates machine-readable. I still wrote a custom parser for all the cases where the marker (geo) is not produced.

Some articles do not have coordinates, although they clearly describe an event, e.g. *Maspero demonstrations*⁵⁹. From articles like these, location candidates are scraped and then resolved:

Listing 6: Article’s locate algorithm

```

1 function locate(article):
2     if 'geo' in article.HTML: # checking for machine-readable coords
3         location = custom.parse(article.geo.text)
4         return
5     else: # look for location candidates
6         candidates = []
7         # all links from the first info box's location field (wikitext)
8         candidates.extend(templates.infoboxes[0].location.links)
9         # all flags in the first info box
10        candidates.extend(templates.infoboxes[0].flags) #HTML
11        # all links from the first paragraph
12        candidates.extend(article.first_paragraph.links) # HTML
13    if len(candidates) and article.isMainArticle:
14        candidates = candidates[:10]
15        until location:
16            location = locate(retrieve(candidates.pop()))
17    return location

```

The above mentioned *Maspero demonstrations* article is exemplary for the candidate list mechanism. Its info box’s location field offers three links to articles of a place: Maspiro, Cairo, Egypt. They are checked until an article with coordinates is found, in this case after the second try, Cairo⁶⁰.

⁵⁹ http://en.wikipedia.org/wiki/Maspero_demonstrations (visited on 01/25/2012)

⁶⁰ <http://en.wikipedia.org/wiki/Cairo> (visited on 01/25/2012)

4.4.4 *Collective authorship*

Most of the authorship processing is being done by SoNet's Toolserver script⁶¹, see [Toolserver](#). The PHP-script directly queries a live copy of the Wikipedia database for all revisions of the requested article. For each revision's author, an edit counter is incremented, and if the author was anonymous, the IP is being resolved to a geographic location. A collection of authors and all resolved locations are then returned in a JSON object.

From the response, my application then creates an author collection. If an author is a registered bot, the author is excluded. From the location list, a second collection is created to manage all author locations for the current article.

The attribution of text passages to authors is done by the web service WikiTrust, see [Third-party sources/Web services](#). It returns an annotated markup that can easily be parsed with the following regular expression (JS):

```
/{{#t:\d+,\d+,[^}]*}}/g;
```

Barring bots, all edits are considered relevant, i.e. reverts or blanking⁶² are not treated in a special way.

4.4.5 *Locating users*

The location of all anonymous authors has already been determined by SoNet's Toolserver script. For each of the remaining (registered) users, it is first checked if the user is included in the *Poor man's checkuser* list (see [Third-party sources/Web services](#)). In that case, the username can be resolved to an IP address which, in turn, can be resolved to an accurate location, see [IP Look-up](#). For all remaining authors, their user page is parsed for a location, see [Parsing user pages](#).

4.4.5.1 *IP Look-up*

Resolving the location⁶³ for a given IP address is a simple call to Quova's IP-lookup API.⁶⁴ The limits of the non-commercial license — 2 requests per second and a maximum of 1,000 requests per day — have been overcome by a caching server proxy.

⁶¹ <https://github.com/volpino/toolserver-scripts/blob/master/\ac{PHP}/api.\ac{PHP}> (visited on 01/25/2012)

⁶² The illegitimate removal of all content of an article, see http://en.wikipedia.org/wiki/Wikipedia:VANDTYPES#Types_of_vandalism (visited on 01/25/2012)

⁶³ Regarding their accuracy and coverage, their website is rather vague, but coverage is at least 99.8% on a state level, see <http://www.quova.com/what/> (visited on 01/25/2012)

⁶⁴ For the API call, a registered account is needed. This provides a secret key which has to be used to sign each request.

4.4.5.2 Parsing user pages

A user page is scanned for possible locations by parsing its HTML content. This is easier than parsing the wikitext, since embedded templates can be inconsistent⁶⁵. The parser looks for links with country names and then checks if they appear in a certain context:

Listing 7: User page location algorithm

```

1 candidates = []
2 for link in userpage.links:
3     if is_country(link.title):
4         candidates.append(link)
5
6 patterns = [" comes? from", " am from", "This user is from",
7            "This user is in", " lives? in", " currently living in"]
8
9 for candidate in candidates:
10     context = link.parent
11     for pattern in patterns:
12         if context.match(pattern):
13             country = candidate
14             break
15 return country

```

Plainly parsing for locations and flags, as done in the article's location parsing, yields too many false positives as some user pages are flooded with flags of countries the user professes to have visited. However, looking at the context of where a link appears, helps in parsing prose such as:

“...and live in a rather small town close to
my country's capital”.

from the user page of *Nightstallion*⁶⁶.

4.4.5.3 Geographic resolution

As accurate as the coordinates from the IP location services may appear, adding the results of the user page analysis to the set of located authors means having to settle for a country-level resolution.

Both the article location and the user page parsing algorithms are searching for countries. The location for a country is looked up in the application's country list (see COUNTRIES in [Wikipedia website](#)). There, each country name is mapped to the geographic coordinates present in its Wikipedia article. For bigger countries, the coordinates refer to the location of the capital, e.g. *United States*⁶⁷ while for smaller ones a central point is denominated,

⁶⁵ See http://en.wikipedia.org/wiki/Category:Nation_of_origin_user_templates (visited on 01/25/2012)

⁶⁶ <http://en.wikipedia.org/wiki/User:Nightstallion> (visited on 01/25/2012)

⁶⁷ http://en.wikipedia.org/wiki/United_States (visited on 01/25/2012)

e.g. Bahrain⁶⁸. In any case, the coordinates still refer to a point on the globe and, by definition, cannot describe an area. Despite this shortcoming, coordinates are the primary means to reference places in Wikipedia.

As a result, all users that were located by user page parsing, are referenced with coordinates of a country. This has the potential to grossly bias seemingly accurate metrics like the signature distance, described in the next section. For a political analysis, however, it may be less important that all user that have been located in the US by user page parsing, share the same coordinates.

4.4.6 Signature distance

In his dissertation, Hardy [24] developed a proximity metric called *signature distance*. This calculates the “average distance between an article and all its contributing authors, weighted by the relative work per author”⁶⁹.

Hardy’s formula for the signature distance uses basic properties of articles and authors. Using the same notation, let ρ be an author and α be an article. For a sample of articles S let A be the set all articles in the sample and P be the set of all located authors in the sample:

$$A = \{\alpha : \alpha \in S\} \quad P = \{\rho : \rho \in S\}$$

Then, $\eta(\rho, \alpha)$ are the contribution(s) of author ρ to article α and $N(\alpha)$ are all contribution(s) to article α :

$$N(\alpha) = \{\eta(\rho, \alpha) : \rho \in P\}$$

Conversely, $P(\alpha)$ are the author(s) who have made contributions to article α :

$$P(\alpha) = \{\rho : \rho \in N(\alpha)\}$$

To calculate the relative work, let $w(\rho, \alpha)$ be relative edit frequency for author ρ on article α :

$$w(\rho, \alpha) = \frac{|\eta(\rho, \alpha)|}{|N(\alpha)|}$$

Let further $\delta(\rho, \alpha)$ be the geodesic distance between author ρ and article α . Then, averaging the distances of authors weighted by their relative work leads to the signature distance $D(\alpha)$:

$$\begin{aligned} D(\alpha) &= \sum_{\rho \in P(\alpha)} \frac{|\eta(\rho, \alpha)| \cdot \delta(\rho, \alpha)}{|N(\alpha)|} \\ &= \sum_{\rho \in P(\alpha)} (w(\rho, \alpha) \cdot \delta(\rho, \alpha)) \end{aligned} \quad (4.1)$$

⁶⁸ <http://en.wikipedia.org/wiki/Bahrain> (visited on 01/25/2012)

⁶⁹ See p. 52 for the complete development of the formula. For clarity, this quotation has been stripped of mathematical symbols.

Hardy weighted the author distances by their edit frequency and thereby treated all contributions as equal. With more data available, the namely attribution of text to individual authors (see [Collective authorship](#)), I can provide two variations of the signature distance metric. Given a sequence of revisions, these metrics provide a distance metric that was calculated using only those revisions that produced the text a visitor sees when reading the article.

4.4.6.1 Edit survival

Following the previously used notation, I will develop the *e.surv signature distance* metric. For a given revision, it counts only those contributions that have provided text still present in that revision. Let Π be the set of all revisions in the sample S :

$$\Pi = \{\pi : \pi \in S\}$$

Let $\epsilon(\rho, \alpha, \pi)$ denote the contribution(s) of located author ρ to article α that are still present in revision π . $E(\alpha, \pi)$ are all contribution(s) to article α that are still present in revision π :

$$E(\alpha, \pi) = \{\epsilon(\rho, \alpha, \pi) : \rho \in P\}$$

Thus the modified weight function $w_{e.surv}$:

$$w_{e.surv}(\rho, \alpha, \pi) = \frac{|\epsilon(\rho, \alpha, \pi)|}{|E(\alpha, \pi)|}$$

This leads to the survivor signature distance :

$$\begin{aligned} D_{e.surv}(\alpha, \pi) &= \sum_{\forall \rho \in P(\alpha)} \frac{|\epsilon(\rho, \alpha, \pi)| \cdot \delta(\rho, \alpha)}{|E(\alpha, \pi)|} \\ &= \sum_{\forall \rho \in P(\alpha)} (w_{e.surv}(\rho, \alpha, \pi) \cdot \delta(\rho, \alpha)) \end{aligned} \tag{4.2}$$

4.4.6.2 Text survival

The second variation, the *t.surv signature distance* metric, makes one further adjustment. It weighs the contributions by the length of their text that survived. For a given revision, it counts how much text can be attributed to which contribution.

Let $\tau(\rho, \alpha, \pi)$ denote the text of contribution(s) of located author ρ to article α that are still present in revision π . $T(\alpha, \pi)$ is all text from contribution(s) to article α that are still present in revision π :

$$T(\alpha, \pi) = \{\tau(\rho, \alpha, \pi) : \rho \in P\}$$

In other words, T is the part of the text in a revision that could be located. The modified weight function then determines an author's share of the located text $w_{t.surv}$:

$$w_{t.surv}(\rho, \alpha, \pi) = \frac{|\tau(\rho, \alpha, \pi)|}{|T(\alpha, \pi)|}$$

This leads to the $t.surv$ signature distance :

$$\begin{aligned} D_{t.surv}(\alpha, \pi) &= \sum_{\forall \rho \in P(\alpha)} \frac{|\tau(\rho, \alpha, \pi)| \cdot \delta(\rho, \alpha)}{|T(\alpha, \pi)|} \\ &= \sum_{\forall \rho \in P(\alpha)} (w_{t.surv}(\rho, \alpha, \pi) \cdot \delta(\rho, \alpha)) \end{aligned} \quad (4.3)$$

4.4.6.3 Implementation of the signature distance

Given that an article has a location and its authors have been located, their signature distance (equation 4.1) can be calculated using the following implementation:

Listing 8: Signature distance algorithm

```

1 sd = 0 # signature distance
2 total = 0 # number of all edits
3 for author in article.authors:
4     if loc in author: # count only authors that have a location
5         dist = geodesic(author.loc, article.loc)
6         edits = author.count # work is the edit count
7         total += edits # postpone relativization to last line
8         sd += dist * edits
9 return sd / total

```

The algorithm in listing 8 relies on the preprocessed edit counts by SoNet's Toolserver script and provides the signature distance only for the latest revision. The computation of the signature distance for all revisions is more expensive. Given at least two differently-located authors, the signature distance changes with each new revision that has a locatable author. That means for each revision, all previously located edits have to be counted. As the computation for revision η_n would have to go over the same edits as the computation for η_{n+1} , the algorithm (see listing 9) uses a technique called *memoization* where previous results are stored within the function for later use.

Listing 9: Signature distance algorithm for all revisions

```

1 function compute(i, located):
2     revision = located[i]
3     dist = geodesic(revision.author.loc, article.loc)
4     if i == 0:
5         return dist
6     return (dist + (i - 1) * compute(i - 1, located)) / i;
7
8 memoized = memoize(compute)
9 # only revisions that have a located author

```

```

10 located = revisions.filter_location()
11 sd = 0
12 for revision, index in located:
13     sd = memoized(index, located)
14     revisions.set(sd)

```

The other two variations of the signature distance have been implemented in a similar way.⁷⁰

4.4.7 Localness

An essential part of the authorship analysis deals with the question whether an author is *local*. If the geographic coordinates are available for an author's location, the geodesic distance can be used to determine the proximity. In this case the users are regarded as local if their geodesic distance to the article location does not exceed a limit. The distance has to be in the lower quartile of all author distances and not be more than 500km.

When the user's location could only be determined by parsing the user page, the user is considered local when the parsed country is the same as the country of the event article.

Listing 10: Localness of an author

```

1 function is_local(author, article, authorship):
2     if author.location:
3         if author.location.coordinates:
4             # coordinates found via IP lookup
5             distance = geodesic(author.location, article.location)
6             # quartiles of all author distances
7             distance_distribution = distance_quartiles(authorship)
8             # within lower quartile or hard limit of 500 km
9             return distance <= min(500,
10                 distance_distribution.1st_quartile)
11         else:
12             # country from user pages
13             return author.location.country == article.location.country
14     else:
15         return null # third state, localness is unknown

```

4.5 VISUALIZATION

The algorithmic analyses of the articles are in fact measurements, i.e. multiple series of numbers, that can also be represented in graphical terms. Following the maxim "The purpose of visualization is insight, not pictures." [33, p. 6], the application's diagrams aim to support the presentation of results and to invite exploration.

⁷⁰ Below, in listing 11, another algorithm based on text survival is presented.

4.5.1 Maps

To show the location that was extracted from an article, a simple world map is used. A marker shows the position of the location's geographic coordinates.



Figure 3: Article location for *2011–2012 Bahraini uprising*

The author analysis groups the located authors by country. For each country, the number of edits of its authors is counted, resulting in a mapping of country name to edit count. A choropleth map (see figure see figure 4) is used to show how this measure varies over different countries. The darker a country is rendered, the higher is its edit count.

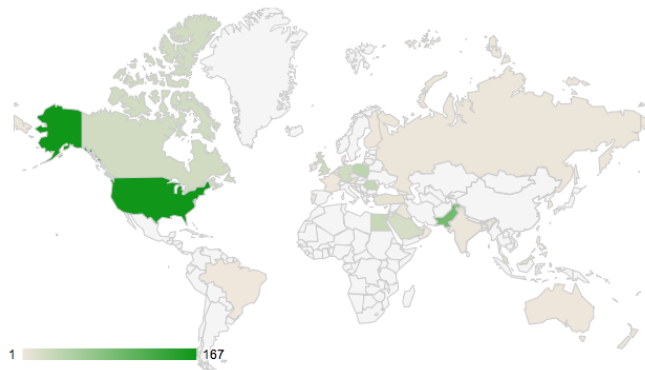


Figure 4: Geographic origins by country for located authors of *2011–2012 Bahraini uprising*

However, this choropleth map only measures the edit counts and does not acknowledge the fact that contributions may disappear by edits in later revisions. Therefore, a second choropleth map (see figure 5) is shown below the first one, displaying text volume based on text survival (see also [Contributions](#)). This allows for a direct comparison of contributor countries with a high activity (edit counts) versus contributor countries whose citizens write text that survives the scrutiny of other editors, e.g. Egypt gained intensity (darker green) in the second map.

For the second choropleth map, the annotated markup provided by WikiTrust is parsed to extract the authors that introduced the text sequences that make up the final text. This sub-




Figure 5: Text survival in revision 471577075 grouped by country for located text of 2011–2012 Bahraini uprising

set of revisions is again grouped by the country of its author (if located):

Listing 11: Edit weight for map

```

1 revision = article.revisions.last
2 total = revision.length
3 countries = {} # country name -> [length of sequence, ...]
4 edits = {} # country name -> proportion of whole text
5 for author in revision.authors:
6     if loc in author: # add only located authors
7         if loc.country not in countries:
8             countries[loc.country] = []
9             countries[loc.country].append(author.text.length)
10 for country in countries.keys:
11     edits[country] = sum(countries[country]) / total
12 return edits

```

4.5.2 Line charts

A timeline chart is used to display the evolution of the several metrics over time: the signature distance, the distance of single revision, and the number of page views for that day, see figure 6. A peak in the orange line (page views) could signal elevated interest in the content while a tight zig-zag pattern in the red line (revision distance) could be an edit war between local and distant contributors.




Figure 6: Activity chart for 2011–2012 Bahraini uprising

4.5.3 Motion chart

The motion chart offers an alternative view on the metrics shown in the choropleth map. This chart type is ideal to follow the change in several indicators over time. Each country is represented by a bubble that, depending on the metric chosen, either moves along an axis or changes its size.⁷¹ The application uses Google's implementation of the motion chart⁷². It is interactive and invites the user to explore the data by allowing the user to freely choose which metric should be represented by which axis.




Figure 7: Temporal development of edit counts by country for 2011–2012 Bahraini uprising

Even clearer than the choropleth map, the motion chart identifies the main contributor countries, e.g. the big blue bubble (Bahrain) and the red one (United States), see figure 7. In the example, changing the metric on the y-axis to show the proportion of located text reveals that contributions from Pakistan have a higher survival rate (the yellow bubble moved up), see figure 8.

4.6 HYPOTHESIS ANALYSIS

The data gathered in the article analysis will be used to test the hypotheses. In addition to the basic requirements (see [Article requirements](#)) for each article to qualify, the following sections describe for each hypothesis which data from the content analysis is used in testing support for it.

For hypotheses that rely on a correlation for their support, the Pearson product-moment correlation coefficient is used. *outliers*

⁷¹ The list of countries has to be limited to the top 10 countries for each measure. Otherwise, the amount of the data would have been too big for a browser application; the worst case, an 11-year-old article with daily contributions from each country, yields more than 1m values per measure.

⁷² <http://code.google.com/apis/chart/interactive/docs/gallery/motionchart.html> (visited on 01/25/2012)




Figure 8: Temporal development of text proportion by country for 2011–2012 Bahraini uprising

H1: Articles are created with only a short delay after the start date of the event.

When the date parsing algorithm finds a date, an article is assumed to treat an event. Moreover, the parser supports different time resolutions. Some events have a precise date-to-date interval, e.g. *2011 Dohuk riots*⁷³ while others merely happened over the course of a month, e.g. *February 2010 Australian cyberattacks*⁷⁴, or even years, e.g. *Oyster Wars*⁷⁵.

An article qualifies for this hypothesis if the date parsing resulted in dates with a day or month resolution. The hypothesis is supported when articles were created within a certain time after the event's start date. For an article start date with a day resolution this limit is 7 days; for articles with a month resolution the limit is 30 days.

H2: The more recent an article, the shorter is the delay between the event start and article creation.

The qualification criteria are the same as for H1. For each article the time difference between article creation and event start date is calculated. The hypothesis is supported if there is a linear dependence between article creation date and the time difference, in the form of a negative correlation coefficient.

H3: Articles are being created first in the English Wikipedia.

Only articles that exist in multiple language editions qualify. Further, the event's location parsing must yield a country where

⁷³ http://en.wikipedia.org/wiki/2011_Dohuk_riots (visited on 01/27/2012)

⁷⁴ http://en.wikipedia.org/wiki/February_2010_Australian_cyberattacks (visited on 01/27/2012)

⁷⁵ http://en.wikipedia.org/wiki/Oyster_Wars (visited on 01/29/2012)

English is not an official language. Then, articles lend support to the hypothesis if they have been created first in the English Wikipedia.

The language codes used to identify Wikipedia's language editions "mostly correspond to the language codes defined by ISO 639-1".[34] This allows to test whether a language edition corresponds to one of a country's official languages (see OFFICIAL LANGUAGES in [Wikipedia website](#)).

H4: Articles about political events are created by people in the events' proximity.

Articles qualify if their creator has been located. An article lends support to the hypothesis if the creator is considered local, see [Localness](#).

H5: In the beginning of the event anonymous users contribute more than registered users.

Articles qualify that have at least 10 revisions (excluding bots) in the time interval after creation. The interval for event articles with day resolution is 7 days; the limit for articles with a month resolution is 30 days. An article lends support to the hypothesis when the number of anonymous contribution is bigger than the number of contributions done by registered users.

H6: For the duration of the event the majority of contributions are local.

Articles qualify when they have at least 10 revisions (where the author was located, excluding bots) for the duration of the event. An article lends support to the hypothesis when the number of local contribution is bigger than the number of distant, i.e. non-local, contributions (see [Localness](#)).

H7: Articles of a political event that has ended will continuously shrink in size.

Articles qualify that have an event end date before the date they were analyzed and have at least 10 revisions after their end date. The hypothesis is supported if there is a linear dependence between the time passed after the end date and the article size, in the form of a negative correlation coefficient.

H8: After an event has ended, there will be more contributions from registered users than from anonymous ones.

Articles qualify that have an event end date before the date they were analyzed and at least 10 revisions (excluding bots) in the time after the event has ended. An article lends support to the hypothesis when the number of anonymous contribution is bigger than the number of contributions done by registered users.

H9: After an event has ended, the spatial distribution of the contributors will become less local.

Articles qualify that have an event end date before the date they were analyzed and at least 10 located revisions (excluding bots) in the time after the event has ended. A shift in the spatial distribution from distant to local coincides with a decrease in signature distance. The hypothesis is supported if there is a linear dependence between the time passed after the end date and the signature distance. If overall the signature distance rises over time, the correlation coefficient will be positive.

correlation enough?

H10: For the duration of the event the majority of located article content comes from local contributions.

Articles qualify that have at least 10 located revisions during the event interval. An article lends support to the hypothesis when local (see H6) contributions during the event interval outnumber the non-local ones (see [Localness](#)).

H11: After an event has ended, the spatial distribution of the surviving contributions will become less local.

Articles qualify that have an event end date before the date they were analyzed and at least 10 revisions from after the time the event has ended and for which a text survival analysis has been done. Similar to H9, this hypothesis is supported if there is a linear dependence between the time passed after the end date and the revisions' signature distance. *make this text ratio dependent* If overall the signature distance rises over time, the correlation coefficient will be positive. *correlation enough?*

4.7 POSSIBLE ENHANCEMENTS

The previously mentioned algorithms can be improved in various ways. In this section I'll suggest starting points for enhancements.

4.7.1 *Edit relevance*

Improvements regarding edit relevance aim to filter revisions that do not seem to change the meaning of an article. Barring revisions coming from bots, all revisions are being treated as relevant. It may be worth investigating whether revisions marked as minor edits⁷⁶ should be counted as full contributions. Also, detecting vandalism could help in filtering for relevant edits, although it can be argued that vandalism is just another form of contribution, albeit a rather drastic one.

4.7.2 *User page parsing*

Parsing the user pages is a problem of information extraction. Currently, the pages are scanned for a fixed number of patterns, e.g. "lives in...". The list of patterns could be extended and tested against user pages. One candidate for this are the WikiProject⁷⁷ info boxes which registered users can put on their user pages. From a box like *Template:WikiProject New York City*⁷⁸ an affiliation with New York City could be inferred.

The current algorithm is scanning for links to Wikipedia and then checks if they appear in the right context. For prose that has not been formatted with links, e.g. "I am from Warsaw." a web service like WIKIPEDIAMINER⁷⁹ could be used.

4.7.3 *Geographic profiling*

Lieberman and Lin [26] describes an algorithm to determine the location of a registered author based on the pages the author contributes to. An implementation would load a user's contribution (MediaWiki API query USERCONTRIBS), fetch these pages and scan for coordinates, construct a convex hull around them and then determine the center.

⁷⁶ See [Revision history](#).

⁷⁷ <http://en.wikipedia.org/wiki/Wikipedia:WikiProject> (visited on 01/27/2012)

⁷⁸ http://en.wikipedia.org/wiki/Template:WikiProject_New_York_City (visited on 01/27/2012)

⁷⁹ Website: <http://wikipedia-miner.cms.waikato.ac.nz/> (visited on 01/27/2012), Example "I am from Warsaw.": <http://wikipedia-miner.cms.waikato.ac.nz/demos/annotate/?source=I+am+from+Warsaw.> (visited on 01/27/2012)

5

EXPERIMENTS

5.1 DATA SETS

- *Using Wikipedia's category system and template embedders, link to foundation sections*
- *Choosing the "right" category*
- *Revolutions by country¹*
- *Can it be representative?*
- *Articles are categorized by people.*

5.1.1 *By category*

5.1.2 *By template*

5.1.3 *Political article vs. place article*

link to appendix with complete list

5.2 APPLICATION RUN

describe one run, time it takes, requests made, small article, big article, images

5.2.1 *Distribution*

5.2.2 *Text survival*

5.2.3 *Localness*

5.2.4 *Motion chart*

- *some examples on accuracy for different countries*
- *clustering of origins: areas of influence*

¹ http://en.wikipedia.org/wiki/Category:Revolutions_by_country

Part III

RESULTS

6

RESULTS

- *for each set, show hypothesis results, values*
- *box charts*

CONCLUSION

- *summarize method*
- *summarize results*
- *which hypotheses got confirmed?*
- *wikipedia as news medium vs history book*

7.1 LIMITATIONS

- *Mobile contributions, smartphones*
- *Privacy*
- *Active prevention by proxies and anonymizers:*
J.A. Muir and P.C. van Oorschot. Internet geolocation and evasion. Tech. rep. Citeseer, 2006
J.A. Muir and P.C.V. Oorschot. „Internet geolocation: Evasion and counterevasion.“ In: ACM Computing Surveys (CSUR) 42.1 (2009), p. 4
M. Duckham and L. Kulik. „A formal model of obfuscation and negotiation for location privacy.“ In: Pervasive Computing (2005), pp. 152–170
- *IE approach with Machine Learning* L. Xiao et al. „Information extraction from the web: System and techniques.“ In: Applied Intelligence 21.2 (2004), pp. 195–224
- *unsupervised IE:* O. Etzioni et al. „Unsupervised named-entity extraction from the web: An experimental study.“ In: Artificial Intelligence 165.1 (2005), pp. 91–134

7.1.1 Political events

The same hypotheses may be applicable to other types of articles than political ones. The key requirements are that they have a location attribute and a time interval. This is easily fulfilled by disaster articles, e.g. Fukushima Daiichi nuclear disaster¹.

¹ http://en.wikipedia.org/wiki/Fukushima_Daiichi_nuclear_disaster (visited on 10/31/2011)

7.1.2 *Article location*

Although *location* is central to more abstract concepts like *Culture*² these subjects clearly defy being attributed with *a* location. Nevertheless, an analysis of the spatial distribution of contributors could be interesting.

7.1.3 *Cross-language article growth*

The growth rates of articles covering the same topic across various language editions could be analyzed to further investigate issues like language barrier—locals contributing only in their language—and information arbitrage as suggested by Adar, Skinner, and Weld [28].

7.2 FURTHER RESEARCH

² <http://en.wikipedia.org/wiki/Culture> (visited on 10/31/2011)

Part IV

APPENDIX

BIBLIOGRAPHY

- [1] The Economist. *The people formerly known as the audience*. URL: <http://www.economist.com/node/18904124> (visited on 08/10/2011).
- [2] The Economist. *Protest in Egypt: Another Arab regime under threat*. 2011. URL: <http://www.economist.com/node/18013760>.
- [3] sueddeutsche.de. *Krise in Ägypten - Die Kinder des 6. April und der Tag der Entscheidung - Politik*. 2011. URL: <http://www.sueddeutsche.de/politik/krise-in-aegypten-die-kinder-des-april-rufen-zum-protest-1.1053426>.
- [4] Wikipedia. *Revision history of 2011 Egyptian revolution*. URL: http://en.wikipedia.org/w/index.php?title=2011_Egyptian_revolution&dir=prev&action=history.
- [5] J. Giles. „Internet encyclopaedias go head to head.“ In: *Nature* 438.7070 (2005), pp. 900–901. ISSN: 0028-0836.
- [6] Wikipedia. *Wikipedia Page Views*. URL: <http://stats.wikimedia.org/EN/TablesPageViewsMonthly.htm> (visited on 08/10/2011).
- [7] A. Chadwick. *Routledge handbook of Internet politics*. Taylor & Francis, 2009. ISBN: 0203962540.
- [8] The Economist. *Libya: A civil war beckons*. 2011. URL: <http://www.economist.com/node/18290470>.
- [9] B. Suh et al. „Us vs. them: Understanding social dynamics in Wikipedia with revert graph visualizations.“ In: *Visual Analytics Science and Technology, 2007. VAST 2007. IEEE Symposium on*. IEEE. 2007, pp. 163–170.
- [10] F.Å. Nielsen. „Wikipedia research and tools: Review and comments.“ In: (2011).
- [11] Wikipedia. *Statistics*. URL: <http://stats.wikimedia.org/EN/Sitemap.htm> (visited on 08/10/2011).
- [12] Wikipedia. *History of Wikipedia*. URL: http://en.wikipedia.org/wiki/History_of_Wikipedia (visited on 12/10/2011).
- [13] Wikimedia Foundation. URL: http://en.wikipedia.org/wiki/Wikimedia_Foundation (visited on 01/24/2012).
- [14] Wikimedia Foundation *annual report 2010-2011*. URL: https://upload.wikimedia.org/wikipedia/commons/4/48/WMF_AR11_SHIP_spreads_15dec11_72dpi.pdf (visited on 01/24/2012).

- [15] Wikipedia. *Protection policy*. URL: http://en.wikipedia.org/wiki/Wikipedia:Protection_policy (visited on 11/16/2011).
- [16] *Minor edit*. URL: http://en.wikipedia.org/wiki/Help:Minor_edit (visited on 01/27/2012).
- [17] Wikipedia. *Why create an account?* URL: http://en.wikipedia.org/wiki/Wikipedia:Why_create_an_account%3F (visited on 08/10/2011).
- [18] Fabian Kaelin. *Research:Anonymous edits*. URL: http://meta.wikimedia.org/wiki/Research:Anonymous_edits (visited on 12/10/2011).
- [19] *Wikipedia Bots*. URL: <http://en.wikipedia.org/wiki/Wikipedia:Bots> (visited on 12/10/2011).
- [20] F.B. Viégas et al. „Talk Before You Type: Coordination in Wikipedia.“ In: *Proceedings of HICSS*. Vol. 40. 2007.
- [21] M. Kramer, A. Gregorowicz, and B. Iyer. „Wiki trust metrics based on phrasal analysis.“ In: *Proceedings of the 4th International Symposium on Wikis*. ACM. 2008, pp. 1–10.
- [22] B.T. Adler and L. De Alfaro. „A content-driven reputation system for the Wikipedia.“ In: *Proceedings of the 16th international conference on World Wide Web*. ACM. 2007, pp. 261–270.
- [23] B.T. Adler et al. „Assigning trust to wikipedia content.“ In: *Proceedings of the 4th International Symposium on Wikis*. ACM. 2008, pp. 1–12.
- [24] D. Hardy. „Volunteered geographic information in Wikipedia.“ PhD thesis. UNIVERSITY OF CALIFORNIA, SANTA BARBARA, 2011.
- [25] J.A. Muir and P.C.V. Oorschot. „Internet geolocation: Evasion and counterevasion.“ In: *ACM Computing Surveys (CSUR)* 42.1 (2009), p. 4.
- [26] M.D. Lieberman and J. Lin. „You are where you edit: Locating Wikipedia users through edit histories.“ In: *ICWSM’09* (2009), pp. 106–113.
- [27] D.K. Lewis. *Philosophical papers*. Vol. 2. Oxford University Press, USA, 1987.
- [28] E. Adar, M. Skinner, and D.S. Weld. „Information arbitrage across multi-lingual Wikipedia.“ In: *Proceedings of the second ACM international conference on Web search and data mining*. ACM. 2009, pp. 94–103.

- [29] B. Hecht and D. Gergle. „The Tower of Babel meets Web 2.0: User-generated content and its applications in a multi-lingual context.“ In: *Proceedings of the 28th international conference on Human factors in computing systems*. ACM. 2010, pp. 291–300.
- [30] A. Kittur, E.H. Chi, and B. Suh. „What’s in Wikipedia?: mapping topics and conflict using socially annotated category structure.“ In: *Proceedings of the 27th international conference on Human factors in computing systems*. ACM. 2009, pp. 1509–1512.
- [31] Fernanda B. Viégas, Martin Wattenberg, and Kushal Dave. „Studying cooperation and conflict between authors with history flow visualizations.“ In: *Proceedings of the SIGCHI conference on Human factors in computing systems*. CHI ’04. Vienna, Austria: ACM, 2004, pp. 575–582. ISBN: 1-58113-702-8. DOI: <http://doi.acm.org/10.1145/985692.985765>.
- [32] A. Kittur et al. „Power of the few vs. wisdom of the crowd: Wikipedia and the rise of the bourgeoisie.“ In: *World Wide Web 1.2* (2007), p. 19.
- [33] S.K. Card, J.D. Mackinlay, and B. Shneiderman. *Readings in information visualization: using vision to think*. Morgan Kaufmann, 1999.
- [34] *List of Wikipedias*. URL: http://en.wikipedia.org/wiki/List_of_Wikipedias (visited on 01/30/2012).
- [35] J.A. Muir and PC van Oorschot. *Internet geolocation and evasion*. Tech. rep. Citeseer, 2006.
- [36] M. Duckham and L. Kulik. „A formal model of obfuscation and negotiation for location privacy.“ In: *Pervasive Computing* (2005), pp. 152–170.
- [37] L. Xiao et al. „Information extraction from the web: System and techniques.“ In: *Applied Intelligence* 21.2 (2004), pp. 195–224.
- [38] O. Etzioni et al. „Unsupervised named-entity extraction from the web: An experimental study.“ In: *Artificial Intelligence* 165.1 (2005), pp. 91–134.